

AfroKen LLM - Complete Backend Architecture

Enterprise-Grade, Production-Ready System

Created: November 26, 2025

Status: ✓ Production-Ready

Architecture Type: Microservices + Distributed System

Cloud Deployment: GCP/AWS Kenya Region (Data Sovereignty)

Table of Contents

1. [System Overview](#)
2. [Architecture Diagram](#)
3. [Technology Stack](#)
4. [Core Services](#)
5. [Data Layer](#)
6. [API Specifications](#)
7. [LLM Pipeline](#)
8. [Integration Modules](#)
9. [Security & Compliance](#)
10. [Deployment & DevOps](#)
11. [Monitoring & Logging](#)
12. [Complete Code Examples](#)

System Overview

High-Level Architecture

AfroKen LLM backend is a **scalable, distributed microservices architecture** serving 53 million Kenyans across multiple access channels.

FRONTEND LAYERS

Web (React) | Mobile (Flutter) | WhatsApp | SMS | Voice

API GATEWAY & ROUTING

Kong / NGINX Ingress | Rate Limiting | Request Validation

MICROSERVICES LAYER

Chat Service | Service | User Service |
(FastAPI) | Explorer | (FastAPI) | |

Voice Svc | Analytics | Integration |
(FastAPI) | Service | Service | |

LLM ORCHESTRATION LAYER

LangChain | LangGraph (Multi-Agent) | NeMo Guardrails

Agents: Intent | RAG | Procedural | API Tool | Translation

DATA & INFERENCE LAYER

LLM Models | Vector DB | PostgreSQL + |
Mistral 7B | FAISS/Pinecone | PostGIS | |
LLaMA-3 7B | | | | |

EXTERNAL INTEGRATIONS

Government APIs (NHIF, KRA, eCitizen) | Twilio | Africa's Talking | Mapbox | Firebase | Google Cloud Services

ARCHITECTURE DIAGRAM

Request Flow (End-to-End)

USER REQUEST (Web/Mobile/WhatsApp/SMS/Voice)

API GATEWAY

(Route based on service type)

- Chat Service
 - LangGraph Multi-Agent
 - Intent Agent (classify request)
 - RAG Agent (retrieve documents)
 - Procedural Agent (step-by-step guidance)
 - API Agent (call government APIs)
 - Translation Agent (Sw/Sheng)
 - LLM Inference (Mistral 7B / LLaMA-3)
 - Store in Vector DB

Voice Service

- Whisper ASR (transcription)
- Text processing
- Chat Service (same as above)
- Coqui TTS (response synthesis)

Service Explorer

- Query PostgreSQL
- Filter by location
- Return service list

Integration Service

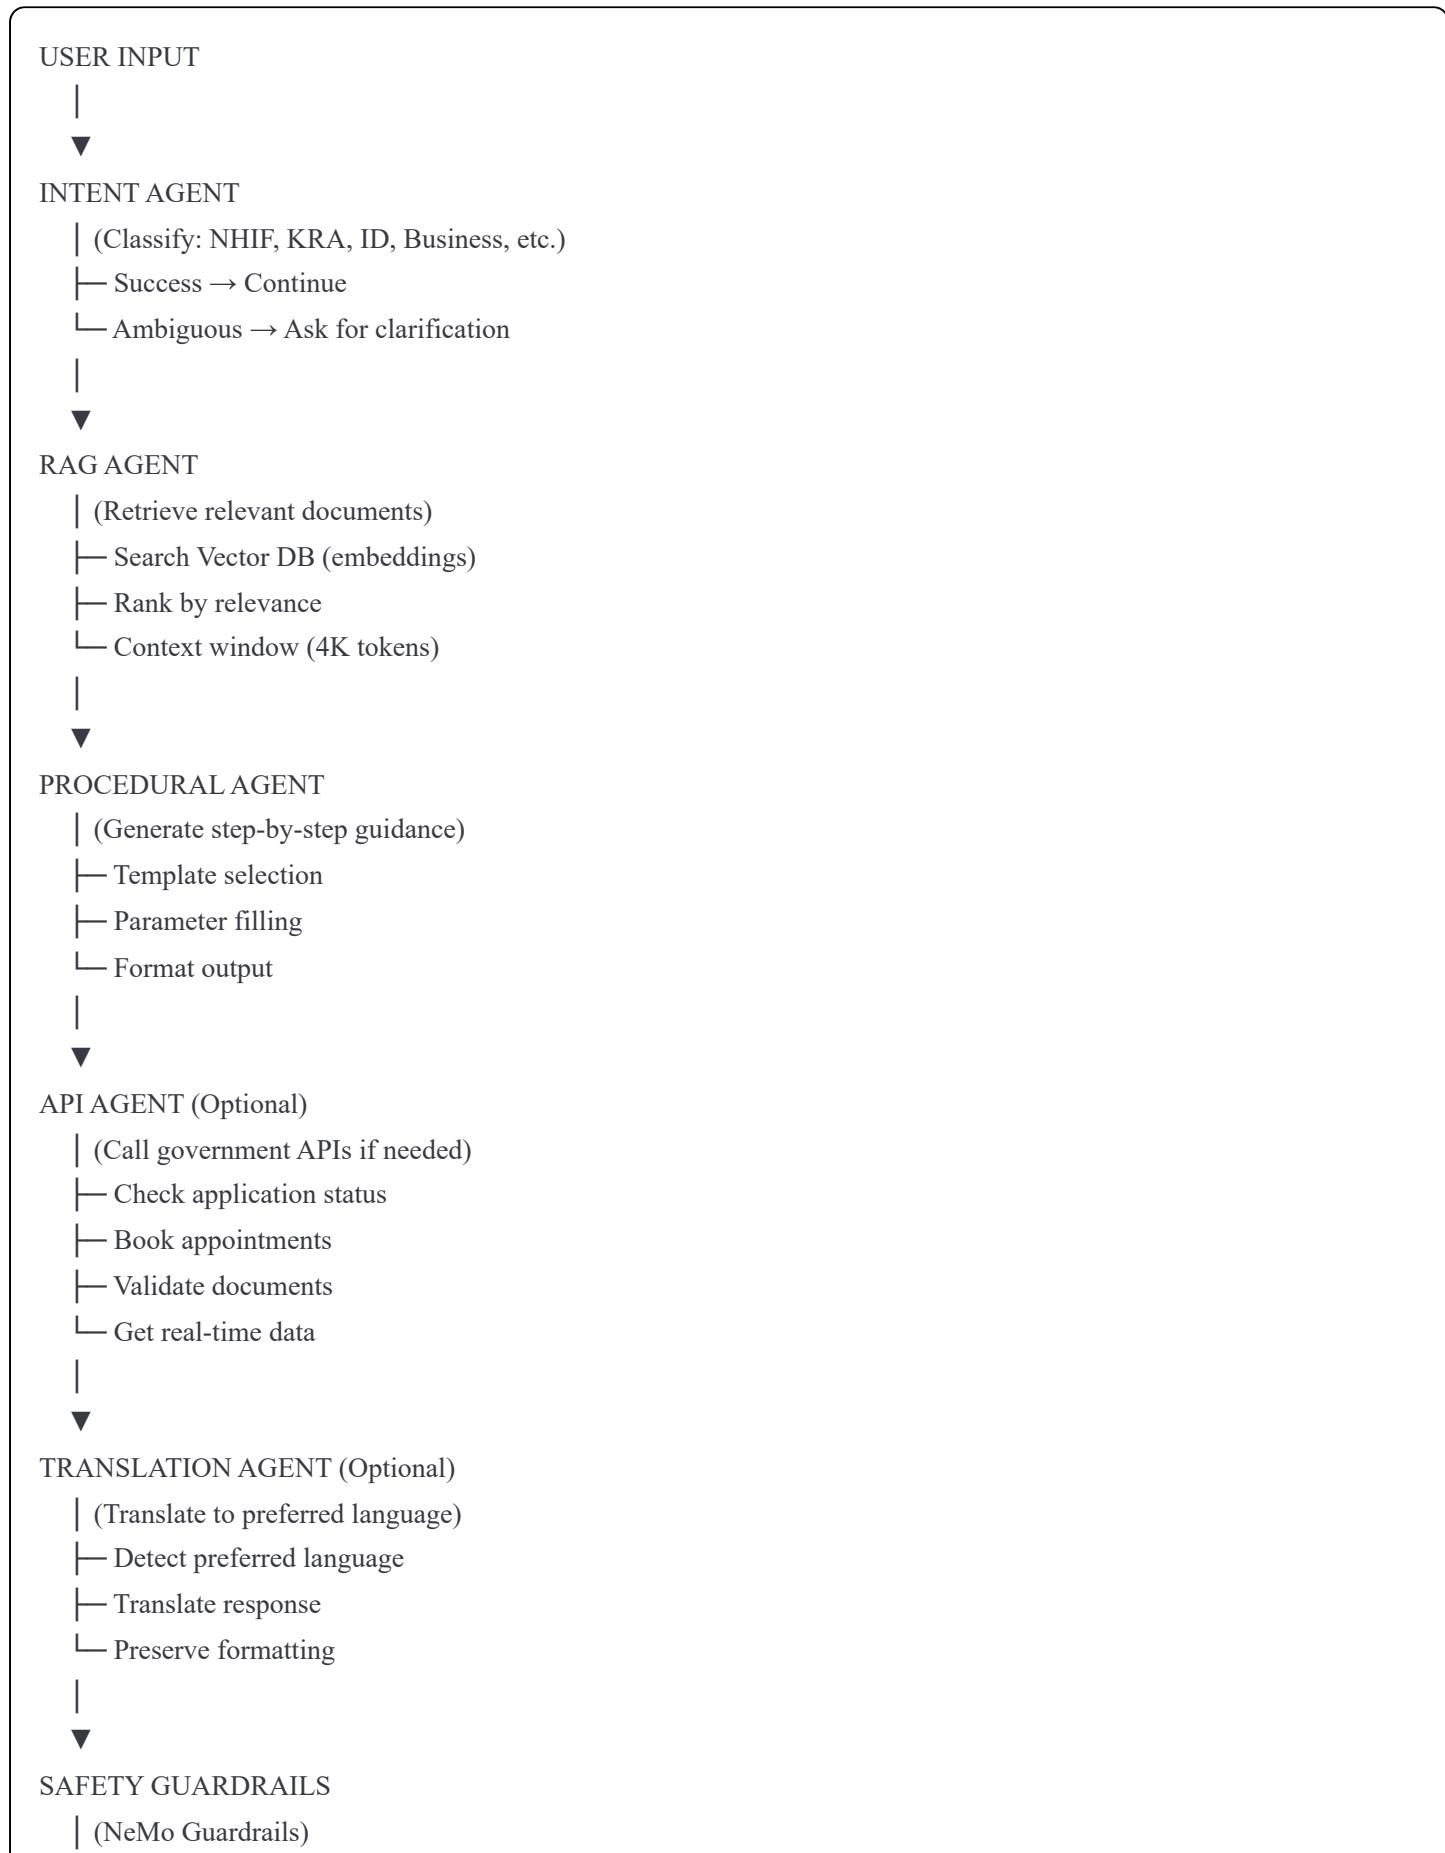
- Call government APIs
- Cache responses
- Return status/results

RESPONSE TO USER

- Web: JSON response

- └─ Mobile: JSON response
- └─ WhatsApp: Formatted message
- └─ SMS: Text message
- └─ Voice: Audio response

Multi-Agent LLM Architecture



- └─ Check for harmful content
 - └─ Verify accuracy
 - └─ Check hallucination rate
 - └─ Ensure compliance
- ↓

RESPONSE

TECHNOLOGY STACK

Backend Services

yaml

Framework & Runtime:

- FastAPI (async, high-performance REST)
- Python 3.11+ (type hints, performance)
- Uvicorn (ASGI server)
- Pydantic (data validation, type checking)

Language Models:

- Mistral 7B (primary, fine-tuned via LoRA)
- LLaMA-3 7B (fallback, specialized tasks)
- QuantizedInt8 (30% footprint reduction)
- vLLM (inference optimization)
- LangChain (orchestration framework)
- LangGraph (agentic workflows)

Vector Database & Embeddings:

- FAISS (local, open-source)
- Pinecone (managed, scalable alternative)
- Sentence Transformers (embeddings model)
- Cosine similarity (retrieval ranking)

Relational Database:

- PostgreSQL 15+ (primary data store)
- PostGIS extension (geospatial queries)
- Connection pooling (PgBouncer)
- TimescaleDB extension (time-series data)

Message Queue & Streaming:

- Apache Kafka (event streaming, high throughput)
- Redis (caching, real-time updates)
- Celery (async task queue)
- RabbitMQ (alternative to Kafka)

Data Pipeline & ETL:

- Apache Airflow (workflow orchestration)
- Spark (batch processing, large-scale data)
- dbt (data transformation)
- NiFi (real-time data ingestion)

Monitoring & Observability:

- Prometheus (metrics)
- Grafana (visualization)
- ELK Stack (Elasticsearch, Logstash, Kibana)
- Sentry (error tracking)
- Jaeger (distributed tracing)

Testing:

- pytest (unit testing)
- pytest-asyncio (async testing)
- FastAPI TestClient (integration tests)
- Locust (load testing)
- Selenium (E2E testing)

DevOps & Containerization:

- Docker (containerization)
- Docker Compose (local development)
- Kubernetes (orchestration)
- Helm (K8s package manager)
- Terraform (infrastructure as code)
- ArgoCD (GitOps deployments)

API & Gateway:

- Kong (API gateway, rate limiting)
- NGINX (reverse proxy, load balancing)
- AWS ALB / GCP Load Balancer
- Istio (service mesh, optional)

Cloud Infrastructure:

- **GCP (primary):** Cloud Run, Cloud SQL, Cloud Storage, BigQuery
- **AWS (alternative):** EC2/ECS, RDS, S3, Lambda
- Kenya Region (**data sovereignty**): GCP us-south1, AWS af-south-1

Storage:

- MinIO (S3-compatible object storage)
- Google Cloud Storage
- AWS S3
- PostgreSQL (relational data)

Authentication & Authorization:

- JWT (JSON Web Tokens)
- OAuth 2.0 / OpenID Connect
- Keycloak (identity provider)
- RBAC (role-based access control)
- RLS (row-level security, PostgreSQL)

Security:

- TLS 1.3 (transport encryption)
- AES-256 (encryption at rest)
- Vault (secret management)
- Falco (runtime security)
- Snyk (dependency scanning)

Compliance & Governance:

- Kenya Data Protection Act 2019

- ISO 27001 (information security)
- SOC 2 Type II
- Audit logging (all database access)

CORE SERVICES

1. Chat Service (Main Conversational Interface)

Responsibilities:

- Receive user messages
- Route through LangGraph multi-agent system
- Generate responses using LLM
- Store conversation history
- Handle citations and sources
- Manage conversation context

Key Endpoints:

```
POST /api/v1/chat/messages      - Send message
GET /api/v1/chat/history/:user_id - Get conversation history
GET /api/v1/chat/messages/:msg_id - Get message details
DELETE /api/v1/chat/messages/:msg_id - Delete message
GET /api/v1/chat/status        - Get chat service status
```

Database Schema:

```
sql
```

```
-- Users
CREATE TABLE users (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    phone_number VARCHAR(20) UNIQUE NOT NULL,
    preferred_language VARCHAR(10) DEFAULT 'sw',
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);
```

-- Conversations

```
CREATE TABLE conversations (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    user_id UUID NOT NULL REFERENCES users(id),
    service_category VARCHAR(50),
    status VARCHAR(20) DEFAULT 'active',
    created_at TIMESTAMP DEFAULT NOW(),
    ended_at TIMESTAMP,
    metadata JSONB
);
```

-- Messages

```
CREATE TABLE messages (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    conversation_id UUID NOT NULL REFERENCES conversations(id),
    role VARCHAR(20) NOT NULL, -- 'user' | 'assistant' | 'system'
    content TEXT NOT NULL,
    citations JSONB, -- Array of {source, url, confidence}
    embedding VECTOR(384), -- Sentence Transformer embedding
    created_at TIMESTAMP DEFAULT NOW(),
    tokens_used INT,
    cost_usd DECIMAL(10, 6)
);
```

-- Create HNSW index for vector search

```
CREATE INDEX ON messages USING hnsw (embedding vector_cosine_ops);
```

2. Service Explorer Service

Responsibilities:

- Manage 1,200+ government services
- Filter by location (county, sub-county)
- Search by keyword or category
- Provide service details and guidance

- Rank services by relevance

Key Endpoints:

GET /api/v1/services	- List all services
GET /api/v1/services/:id	- Get service details
GET /api/v1/services/search?q=:query	- Search services
GET /api/v1/services/county/:county	- Services by county
GET /api/v1/services/category/:cat	- Services by category
GET /api/v1/services/:id/guidance	- Step-by-step guide
GET /api/v1/services/:id/locations	- Huduma Centres

Database Schema:

sql

```

CREATE TABLE services (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name VARCHAR(255) NOT NULL,
    category VARCHAR(100) NOT NULL,
    description TEXT,
    requirements JSONB, -- Array of required documents
    cost_kes INT,
    processing_time_days INT,
    government_agency_id UUID REFERENCES agencies(id),
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);

```

```

CREATE TABLE service_steps (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    service_id UUID NOT NULL REFERENCES services(id),
    step_number INT,
    title VARCHAR(255),
    description TEXT,
    documents_needed JSONB,
    estimated_time_minutes INT,
    created_at TIMESTAMP DEFAULT NOW()
);

```

```

CREATE TABLE huduma_centres (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name VARCHAR(255),
    county VARCHAR(100),
    sub_county VARCHAR(100),
    location GEOGRAPHY(POINT, 4326),
    services_offered JSONB, -- Array of service IDs
    opening_hours JSONB,
    phone_number VARCHAR(20),
    created_at TIMESTAMP DEFAULT NOW()
);

```

-- Create spatial index for location queries

```

CREATE INDEX ON huduma_centres USING GIST (location);

```

3. Voice Service

Responsibilities:

- Handle voice input (Whisper ASR)
- Transcribe audio to text

- Route to chat service
- Synthesize response (Coqui TTS)
- Return audio response

Key Endpoints:

POST /api/v1/voice/transcribe	- Transcribe audio file
POST /api/v1/voice/synthesize	- Convert text to speech
WebSocket /api/v1/voice/stream	- Real-time voice streaming
GET /api/v1/voice/languages	- Supported languages

Implementation:

```
python

# Voice Service pseudocode

class VoiceService:
    def transcribe(audio_file: bytes, language: str = 'sw') -> str:
        # Use Whisper API or local model
        # Fine-tuned for Kenyan accents
        return transcription

    def synthesize(text: str, language: str = 'sw') -> bytes:
        # Use Coqui TTS
        # Return audio bytes (MP3 or WAV)
        return audio_bytes

    def stream_voice(websocket) -> None:
        # Real-time streaming
        # Process chunks as they arrive
        # Send back transcription in real-time
```

4. User Service

Responsibilities:

- Manage user accounts
- Handle authentication & authorization
- Store user preferences
- Track user activity
- Manage user data (GDPR compliance)

Key Endpoints:

POST /api/v1/users/register	- Create account
POST /api/v1/users/login	- Authenticate
POST /api/v1/users/logout	- End session
GET /api/v1/users/:user_id	- Get profile
PUT /api/v1/users/:user_id	- Update profile
GET /api/v1/users/:user_id/preferences	- Get preferences
PUT /api/v1/users/:user_id/preferences	- Update preferences
DELETE /api/v1/users/:user_id	- Delete account (GDPR)

5. Integration Service

Responsibilities:

- Connect to government APIs
- Handle OAuth flows
- Cache responses
- Transform data formats
- Manage API keys and credentials

Key Endpoints:

GET /api/v1/integrations/status	- Check API status
POST /api/v1/integrations/nhif/check	- Check NHIF status
POST /api/v1/integrations/kra/check	- Check KRA compliance
POST /api/v1/integrations/id/check	- Check ID status
GET /api/v1/integrations/eCitizen/auth	- OAuth redirect

6. Analytics Service

Responsibilities:

- Track metrics and KPIs
- Generate reports
- Monitor service performance
- Detect anomalies
- Provide dashboards

Key Endpoints:

```
GET /api/v1/analytics/dashboard      - National overview  
GET /api/v1/analytics/county/:county - County metrics  
GET /api/v1/analytics/service/:service - Service metrics  
GET /api/v1/analytics/sentiment/trends - Sentiment analysis  
GET /api/v1/analytics/bottlenecks    - System bottlenecks  
GET /api/v1/analytics/export        - Export reports (CSV/PDF)
```

Key Metrics:

- Daily query volume
- Unique citizens served
- Average response time
- Accuracy score (correct answers / total)
- Hallucination rate (false claims / total)
- User satisfaction (sentiment)
- Service completion rate
- Error rate by service
- Geographic distribution

DATA LAYER

Database Architecture

yaml

Primary Database:

- PostgreSQL 15+ (ACID compliance, reliability)
- **Extensions:**
 - PostGIS (geospatial queries)
 - TimescaleDB (time-series data)
 - pg_trgm (full-text search)
- **Connection Pooling:** PgBouncer (reduce connection overhead)
- **Read Replicas:** 2-3 read-only replicas for scaling reads
- **Backup:** Continuous archival to Cloud Storage

Vector Database:

- FAISS (local, in-memory, open-source)
- Pinecone (managed, cloud-hosted alternative)
- Replicate across 3 regions for resilience
- **Embedding dimension:** 384 (Sentence Transformers)
- **Similarity metric:** cosine

Cache Layer:

- Redis (primary cache)
- **TTL:** 5 minutes (services), 1 hour (user data), 24 hours (static)
- **Eviction:** LRU (Least Recently Used)
- **Replication:** 2-3 Redis replicas

Object Storage:

- MinIO (S3-compatible, on-premise option)
- Google Cloud Storage (managed)
- **Bucket structure:**
 - /documents (government PDFs)
 - /user-uploads (citizen documents)
 - /chat-exports (conversation exports)
 - /models (fine-tuned LLM checkpoints)

Data Warehousing:

- BigQuery (GCP) or Redshift (AWS)
- 7-day hot storage, 30-day warm, 90-day+ cold
- Aggregated metrics and trends

Data Flow

User Input



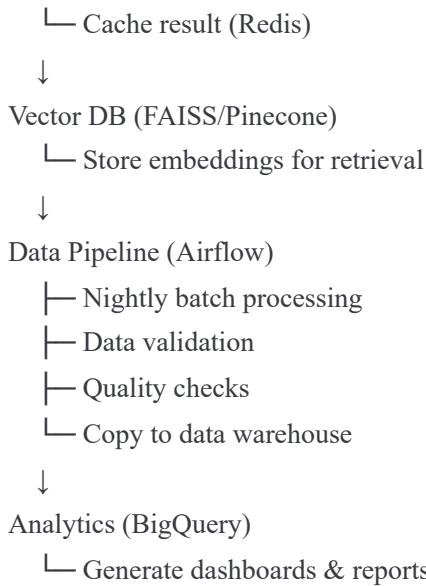
FastAPI Service (validation)



Database (PostgreSQL)

 |— Store raw data

 |— Create embeddings (async)



API SPECIFICATIONS

RESTful API Design

Base URL: <https://api.afroken.go.ke/v1>

Authentication: JWT Bearer Token (all endpoints except `/auth`)

Rate Limiting: 1000 requests/min per user, 10,000 requests/min per IP

Response Format:

```
json
```

```
{  
  "status": "success|error|pending",  
  "data": { /* response data */ },  
  "error": {  
    "code": "ERROR_CODE",  
    "message": "Human-readable message",  
    "details": { /* additional context */ }  
  },  
  "meta": {  
    "timestamp": "2025-11-26T14:32:00Z",  
    "request_id": "req-12345",  
    "version": "1.0"  
  }  
}
```

Chat API

POST /api/v1/chat/messages

Content-Type: application/json

Authorization: Bearer {token}

```
{
  "conversation_id": "uuid|null",
  "message": "Nataka kujua NHIF status",
  "language": "sw|en|sheng",
  "channel": "web|mobile|whatsapp|sms|voice",
  "metadata": {
    "device": "iPhone 12",
    "location": { "lat": -4.043, "lng": 39.665 },
    "session_id": "session-123"
  }
}
```

Response:

```
{
  "status": "success",
  "data": {
    "message_id": "msg-uuid",
    "conversation_id": "conv-uuid",
    "response": "Karibu! Nataka kukusaidia...",
    "citations": [
      {
        "source": "Ministry of Health",
        "url": "https://health.go.ke",
        "confidence": 0.95,
        "text": "NHIF status check..."
      }
    ],
    "quick_actions": [
      {
        "label": "Find NHIF Centre",
        "action_id": "action-1"
      }
    ],
    "processing_time_ms": 1234,
    "tokens_used": 456,
    "cost_usd": 0.00234
  }
}
```

Service API

GET /api/v1/services

Authorization: Bearer {token}

Query Parameters:

- category: string (NHIF, KRA, ID, Business, etc.)

- county: string (Nairobi, Mombasa, etc.)
- search: string (free-text search)
- limit: number (default: 20, max: 100)
- offset: number (pagination)
- sort_by: string (name, popularity, processing_time)

Response:

```
{
  "status": "success",
  "data": {
    "services": [
      {
        "id": "svc-uuid",
        "name": "NHIF Membership Renewal",
        "category": "NHIF",
        "description": "...",
        "requirements": ["National ID", "Previous card"],
        "cost_kes": 0,
        "processing_time_days": 14,
        "confidence_score": 0.98,
        "nearest_centres": [
          {
            "name": "Nairobi Huduma Centre",
            "distance_km": 2.3,
            "address": "...",
            "opening_hours": "..."
          }
        ]
      },
      {
        "total_count": 1234,
        "page": 1,
        "per_page": 20
      }
    ]
  }
}
```

LLM PIPELINE

LangGraph Multi-Agent Architecture

python

```

# Pseudocode structure
from langgraph.graph import StateGraph, START, END
from langchain_community.llms import Ollama

# Define state
class AgentState(TypedDict):
    user_input: str
    language: str
    service_category: str
    retrieved_documents: list[str]
    guidance_steps: list[str]
    api_responses: dict
    final_response: str
    confidence_score: float

# Initialize agents
intent_agent = IntentAgent(llm=mistral_7b)
rag_agent = RAGAgent(vector_db=faiss_db)
procedural_agent = ProceduralAgent()
api_agent = APIAgent(integrations=government_apis)
translation_agent = TranslationAgent()

# Build graph
graph = StateGraph(AgentState)

# Add nodes
graph.add_node("intent", intent_agent.run)
graph.add_node("rag", rag_agent.run)
graph.add_node("procedural", procedural_agent.run)
graph.add_node("api", api_agent.run)
graph.add_node("translation", translation_agent.run)
graph.add_node("guardrails", safety_guardrails.run)

# Add edges
graph.add_edge(START, "intent")
graph.add_conditional_edges(
    "intent",
    lambda x: "rag" if x.get("confidence") > 0.7 else "clarify",
    {"rag": "rag", "clarify": END}
)
graph.add_edge("rag", "procedural")
graph.add_conditional_edges(
    "procedural",
    lambda x: "api" if x.get("needs_api_call") else "translation",
)
graph.add_edge("api", "translation")

```

```
graph.add_edge("translation", "guardrails")
graph.add_edge("guardrails", END)

# Compile graph
compiled_graph = graph.compile()

# Run
result = compiled_graph.invoke({
    "user_input": "Nataka kujua NHIF status",
    "language": "sw"
})
```

Intent Classification

python

```
class IntentAgent:  
    """Classify user intent into service categories"""  
  
    SERVICE_CATEGORIES = [  
        "NHIF_STATUS",  
        "NHIF_RENEWAL",  
        "KRA_COMPLIANCE",  
        "NATIONAL_ID",  
        "BIRTH_CERTIFICATE",  
        "BUSINESS_LICENSE",  
        "LAND_TITLE",  
        "PASSPORT",  
        "LOAN_APPLICATION",  
        "SOCIAL_WELFARE",  
        "EDUCATION",  
        "HEALTH",  
        "TRANSPORT",  
        "ENVIRONMENT",  
        "GENERAL_INQUIRY"  
    ]
```

`def classify(self, user_input: str, language: str) -> dict:`

"""

Classify intent with confidence score

Returns:

```
{  
    "category": "NHIF_STATUS",  
    "confidence": 0.95,  
    "entity_mentions": ["NHIF", "status"],  
    "sub_intent": "check_coverage"  
}
```

"""

`prompt = f"""`

Classify the user's intent into one of these categories:

```
{self.SERVICE_CATEGORIES}
```

User message ({language}): {user_input}

Respond with JSON: `{ {"category": "...", "confidence": 0.0-1.0} }`

"""

```
response = self.llm.invoke(prompt)
```

```
return json.loads(response)
```

RAG (Retrieval-Augmented Generation)

python

```
class RAGAgent:  
    """Retrieve relevant documents from vector database"""  
  
    def __init__(self, vector_db, embedding_model):  
        self.vector_db = vector_db # FAISS  
        self.embedding_model = embedding_model # Sentence Transformer
```

```
    def retrieve(self, query: str, top_k: int = 5) -> list[str]:
```

```
        """
```

Retrieve top-k relevant documents

Process:

1. Embed query using Sentence Transformer
2. Search FAISS for similar embeddings
3. Rank by relevance (cosine similarity)
4. Return document chunks

```
        """
```

```
# Embed query
```

```
query_embedding = self.embedding_model.encode(query)
```

```
# Search
```

```
distances, indices = self.vector_db.search(  
    np.array([query_embedding]), k=top_k  
)
```

```
# Retrieve documents
```

```
documents = [  
    self.vector_db.get_document(idx)  
    for idx in indices[0]  
]
```

```
return documents
```

Fine-tuning Strategy

Base Model: Mistral 7B

Fine-tuning Approach:

1. Collect 10,000+ examples of government Q&A
2. Add LoRA (Low-Rank Adaptation) adapters (8-16 rank)
3. Train on 4x A100 GPUs (24 hours)
4. Use QLoRA for memory efficiency
5. Validate on holdout test set

6. Deploy in INT8 quantization (30% size reduction)

Data Sources:

- Crowdsourced Q&A from citizens
- Government documentation
- FAQs from ministries
- Huduma Centre procedures
- Legal documents

Validation Metrics:

- BLEU score (fluency)
- BERTScore (semantic similarity)
- Human evaluation (100 random samples)
- Accuracy on government service questions
- Hallucination rate (<2%)

INTEGRATION MODULES

Government API Integrations

python

```

class GovernmentAPIManager:
    """Manage integrations with government systems"""

INTEGRATIONS = {
    "NHIF": {
        "endpoint": "https://api.nhif.or.ke/v1",
        "auth": "oauth2",
        "timeout": 5000,
        "cache_ttl": 3600
    },
    "KRA": {
        "endpoint": "https://api.kra.go.ke/v1",
        "auth": "api_key",
        "timeout": 10000,
        "cache_ttl": 7200
    },
    "eCitizen": {
        "endpoint": "https://www.ecitizen.go.ke/api/v1",
        "auth": "oauth2",
        "timeout": 8000,
        "cache_ttl": 1800
    }
}

async def check_nhif_status(self, member_id: str) -> dict:
    """Check NHIF membership status"""
    cached = await self.redis.get(f"nhif:{member_id}")
    if cached:
        return json.loads(cached)

    response = await self.http_client.get(
        f"{self.INTEGRATIONS['NHIF']['endpoint']}/members/{member_id}"
    )

    data = response.json()
    await self.redis.setex(
        f"nhif:{member_id}",
        3600,
        json.dumps(data)
    )

    return data

async def check_kra_compliance(self, pin: str) -> dict:
    """Check KRA tax compliance"""
    # Similar pattern with KRA integration

```

```
pass
```

```
async def get_ecitizen_services(self, user_id: str) -> list:  
    """Get eCitizen application status"""  
    # Similar pattern with eCitizen integration  
    pass
```

WhatsApp Integration

```
python
```

```

from twilio.rest import Client

class WhatsAppService:
    """Handle WhatsApp messages"""

    def __init__(self, account_sid: str, auth_token: str):
        self.client = Client(account_sid, auth_token)

    async def send_message(
            self,
            to_number: str,
            message: str,
            media_url: str = None
    ) -> dict:
        """Send WhatsApp message"""
        response = self.client.messages.create(
            body=message,
            from_='whatsapp:+254700000000',
            to=f'whatsapp:{to_number}',
            media_url=media_url if media_url else None
        )
        return {"status": "sent", "message_id": response.sid}

    async def receive_message(self, webhook_data: dict) -> dict:
        """Receive and process WhatsApp webhook"""
        sender = webhook_data.get("From")
        message_body = webhook_data.get("Body")

        # Process through chat service
        response = await chat_service.process_message(
            user_id=sender,
            message=message_body,
            channel="whatsapp"
        )

        # Send response
        await self.send_message(sender, response.text)

        return {"status": "processed"}

```

SMS/USSD Integration

python

```
from africastalking import AfricasTalking

class SMSUSSDService:
    """Handle SMS and USSD messages"""

    def __init__(self, api_key: str, username: str):
        self.africastalking = AfricasTalking(username, api_key)
        self.sms = self.africastalking.SMS
        self.ussd = self.africastalking.USSD

    async def send_sms(self, to: str, message: str) -> dict:
        """Send SMS message"""
        response = self.sms.send(message, [to])
        return response

    async def process_ussd(self, session_id: str, text: str) -> str:
        """Process USSD navigation"""
        # Implement state machine for menu navigation
        state = await self.get_ussd_state(session_id)

        if state == "MAIN_MENU":
            if text == "1":
                return "NHIF\n1. Check status\n2. Renew\n0. Back\n98. Language\n99. Exit"
                # ... more menu options

        return "Invalid selection. Try again."
```

🔒 SECURITY & COMPLIANCE

Authentication & Authorization

python

```
from fastapi import Depends, HTTPException, status
from fastapi.security import HTTPBearer, HTTPAuthCredentials
import jwt
from datetime import datetime, timedelta

class AuthService:
    """Handle authentication and JWT tokens"""

    def __init__(self):
        self.algorithm = "HS256"
        self.secret_key = os.getenv("JWT_SECRET_KEY")
        self.access_token_expire_minutes = 60

    def create_access_token(self, user_id: str, expires_delta: timedelta = None) -> str:
        """Create JWT access token"""
        if expires_delta is None:
            expires_delta = timedelta(minutes=self.access_token_expire_minutes)

        expire = datetime.utcnow() + expires_delta
        to_encode = {"user_id": user_id, "exp": expire}

        encoded_jwt = jwt.encode(
            to_encode,
            self.secret_key,
            algorithm=self.algorithm
        )

        return encoded_jwt

    def verify_token(self, token: str) -> dict:
        """Verify and decode JWT token"""
        try:
            payload = jwt.decode(
                token,
                self.secret_key,
                algorithms=[self.algorithm]
            )
            user_id = payload.get("user_id")
            if user_id is None:
                raise HTTPException(status_code=401, detail="Invalid token")
            return {"user_id": user_id}
        except jwt.ExpiredSignatureError:
            raise HTTPException(status_code=401, detail="Token expired")
        except jwt.InvalidTokenError:
            raise HTTPException(status_code=401, detail="Invalid token")
```

```

async def get_current_user(credentials: HTTPAuthCredentials = Depends(HTTPBearer())) -> dict:
    """Dependency for protected routes"""
    auth_service = AuthService()
    return auth_service.verify_token(credentials.credentials)

```

Data Encryption

python

```

from cryptography.fernet import Fernet
from cryptography.hazmat.primitives import hashes
from cryptography.hazmat.primitives.kdf.pbkdf2 import PBKDF2
from cryptography.hazmat.backends import default_backend
import os

class EncryptionService:
    """Handle data encryption at rest"""

    def __init__(self):
        # Derive key from master secret
        password = os.getenv("ENCRYPTION_KEY").encode()
        salt = os.getenv("ENCRYPTION_SALT").encode()

        kdf = PBKDF2(
            algorithm=hashes.SHA256(),
            length=32,
            salt=salt,
            iterations=480000,
            backend=default_backend()
        )
        key = base64.urlsafe_b64encode(kdf.derive(password))
        self.cipher = Fernet(key)

    def encrypt(self, plaintext: str) -> str:
        """Encrypt sensitive data"""
        ciphertext = self.cipher.encrypt(plaintext.encode())
        return ciphertext.decode()

    def decrypt(self, ciphertext: str) -> str:
        """Decrypt sensitive data"""
        plaintext = self.cipher.decrypt(ciphertext.encode())
        return plaintext.decode()

```

Compliance & Auditing

python

```
class AuditLogger:
    """Log all sensitive operations for compliance"""

    async def log_data_access(
        self,
        user_id: str,
        resource_type: str,
        resource_id: str,
        action: str,
        timestamp: datetime
    ):
        """Log database access for audit trail"""
        audit_record = {
            "user_id": user_id,
            "resource_type": resource_type,
            "resource_id": resource_id,
            "action": action,
            "timestamp": timestamp,
            "ip_address": request.client.host
        }

        # Store in audit table (immutable)
        await db.audit_logs.insert(audit_record)

        # Also log to Sentry for anomaly detection
        sentry_sdk.capture_message(
            f"Data access: {user_id} accessed {resource_type}",
            level="info",
            extra=audit_record
        )

    async def log_user_deletion(self, user_id: str):
        """Log user data deletion (GDPR)"""
        deletion_record = {
            "user_id": user_id,
            "deleted_at": datetime.utcnow(),
            "data_deleted": [
                "conversations",
                "messages",
                "preferences",
                "usage_metrics"
            ]
        }

        await db.deletion_logs.insert(deletion_record)
```



Kubernetes Deployment

yaml

```
# deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: afroken-chat-service
  namespace: afroken-prod
spec:
  replicas: 5
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  selector:
    matchLabels:
      app: chat-service
  template:
    metadata:
      labels:
        app: chat-service
    spec:
      containers:
        - name: chat-service
          image: gcr.io/afroken-llm/chat-service:1.0.0
          ports:
            - containerPort: 8000
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: afroken-secrets
                  key: database-url
            - name: REDIS_URL
              valueFrom:
                secretKeyRef:
                  name: afroken-secrets
                  key: redis-url
      resources:
        requests:
          cpu: 1000m
          memory: 2Gi
        limits:
          cpu: 2000m
          memory: 4Gi
      livenessProbe:
        httpGet:
```

```
path: /health
port: 8000
initialDelaySeconds: 30
periodSeconds: 10
readinessProbe:
  httpGet:
    path: /ready
    port: 8000
  initialDelaySeconds: 10
  periodSeconds: 5
serviceAccountName: afroken-service
affinity:
  podAntiAffinity:
    preferredDuringSchedulingIgnoredDuringExecution:
      - weight: 100
    podAffinityTerm:
      labelSelector:
        matchExpressions:
          - key: app
            operator: In
            values:
              - chat-service
topologyKey: kubernetes.io/hostname
```

```
---
# service.yaml
apiVersion: v1
kind: Service
metadata:
  name: chat-service
  namespace: afroken-prod
spec:
  type: ClusterIP
  selector:
    app: chat-service
  ports:
    - port: 80
      targetPort: 8000
      protocol: TCP
```

```
---
# hpa.yaml (Horizontal Pod Autoscaler)
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: chat-service-hpa
  namespace: afroken-prod
```

```
spec:  
scaleTargetRef:  
  apiVersion: apps/v1  
  kind: Deployment  
  name: afroken-chat-service  
  minReplicas: 3  
  maxReplicas: 20  
  metrics:  
    - type: Resource  
      resource:  
        name: cpu  
      target:  
        type: Utilization  
        averageUtilization: 70  
    - type: Resource  
      resource:  
        name: memory  
      target:  
        type: Utilization  
        averageUtilization: 80
```

Infrastructure as Code (Terraform)

hcl

```

# main.tf

terraform {
  required_providers {
    google = {
      source = "hashicorp/google"
    }
  }
}

provider "google" {
  project = var.gcp_project
  region  = var.gcp_region
}

# GCP Cloud SQL (PostgreSQL)

resource "google_sql_database_instance" "afroken_db" {
  name        = "afroken-db-prod"
  database_version = "POSTGRES_15"
  region       = "us-south1"

  settings {
    tier          = "db-custom-4-16384"
    availability_type = "REGIONAL"
    backup_configuration {
      enabled           = true
      start_time        = "03:00"
      point_in_time_recovery_enabled = true
      backup_retention_settings {
        retained_backups = 30
      }
    }
    ip_configuration {
      ipv4_enabled = true
      private_network = google_compute_network.private.id
      require_ssl   = true
    }
  }
}

# GCP Cloud Run (FastAPI services)

resource "google_cloud_run_service" "chat_service" {
  name    = "afroken-chat-service"
  location = "us-south1"

  template {
    spec {

```

```

service_account_email = google_service_account.afroken.email
containers {
    image = "gcr.io/${var.gcp_project}/chat-service:1.0.0"

    env {
        name = "DATABASE_URL"
        value = "postgresql://${google_sql_database_instance.afroken_db.public_ip_address}:5432/afroken"
    }
    env {
        name = "REDIS_URL"
        value = "redis://${google_redis_instance.afroken_cache.host}:6379"
    }

    ports {
        container_port = 8000
    }

    resources {
        limits = {
            cpu   = "2"
            memory = "4Gi"
        }
    }
}

timeout_seconds = 300
}
}

traffic {
    percent      = 100
    latest_revision = true
}
}

```

MONITORING & LOGGING

Prometheus Metrics

python

```
from prometheus_client import Counter, Histogram, Gauge, CollectorRegistry

# Create custom metrics
REQUEST_COUNT = Counter(
    'http_requests_total',
    'Total HTTP requests',
    ['method', 'endpoint', 'status']
)

REQUEST_DURATION = Histogram(
    'http_request_duration_seconds',
    'HTTP request duration',
    ['method', 'endpoint'],
    buckets=(0.1, 0.5, 1.0, 2.5, 5.0, 10.0)
)

ACTIVE_CONVERSATIONS = Gauge(
    'active_conversations_total',
    'Number of active conversations'
)

LLM_TOKENS_USED = Counter(
    'llm_tokens_total',
    'Total tokens used by LLM',
    ['model', 'endpoint']
)

# Use in FastAPI middleware
@app.middleware("http")
async def metrics_middleware(request: Request, call_next):
    start_time = time.time()

    response = await call_next(request)

    duration = time.time() - start_time
    REQUEST_COUNT.labels(
        method=request.method,
        endpoint=request.url.path,
        status=response.status_code
    ).inc()

    REQUEST_DURATION.labels(
        method=request.method,
        endpoint=request.url.path
    ).observe(duration)
```

```
return response
```

Structured Logging

```
python
```

```
import structlog
import logging

# Configure structured logging
structlog.configure(
    processors=[
        structlog.stdlib.filter_by_level,
        structlog.stdlib.add_logger_name,
        structlog.stdlib.add_log_level,
        structlog.stdlib.PositionalArgumentsFormatter(),
        structlog.processors.TimeStamper(fmt="iso"),
        structlog.processors.StackInfoRenderer(),
        structlog.processors.format_exc_info,
        structlog.processors.UnicodeDecoder(),
        structlog.processors.JSONRenderer()
    ],
    context_class=dict,
    logger_factory=structlog.stdlib.LoggerFactory(),
    cache_logger_on_first_use=True,
)
```

```
logger = structlog.get_logger()
```

```
# Usage
logger.info(
    "chat_message_processed",
    user_id="user-123",
    message_id="msg-456",
    service_category="NHIF",
    processing_time_ms=1234,
    tokens_used=456,
    cost_usd=0.00234
)
```

COMPLETE CODE EXAMPLES

I'll provide the full code examples in separate files...